

The Rust Way of OS Development

Philipp Oppermann

May 30, 2018 HTWG Konstanz

About Me


- Computer science student at KIT (Karlsruhe)
- *“Writing an OS in Rust”* blog series (os.phil-opp.com)
- Embedded Rust development

Contact:

- phil-opp on GitHub
- Email: hello@phil-opp.com

Rust

- 3 year old programming language
- Memory safety without garbage collection
- Used by Mozilla, Dropbox, Cloudflare, ...



```
enum Event {  
    Load,  
    KeyPress(char),  
    Click { x: i64, y: i64 }  
}  
  
fn print_event(event: Event) {  
    match event {  
        Event::Load => println!("Loaded"),  
        Event::KeyPress(c) => println!("Key {} pressed", c),  
        Event::Click {x, y} => println!("Clicked at x={}, y={}", x, y),  
    }  
}
```

OS Development

- “Bare metal” environment
 - No underlying operating system
 - No processes, threads, files, heap, ...

Goals

- Abstractions
 - For hardware devices (drivers, files, ...)
 - For concurrency (threads, synchronization primitives, ...)
- Isolation (processes, address spaces, ...)
- Security

OS Development in Rust

- **Writing an OS in Rust:** Tutorials for basic functionality
 - Booting, testing, CPU exceptions, page tables
 - No C dependencies
 - Works on Linux, Windows, macOS

Writing an OS in Rust (Second Edition) Philipp Oppermann's blog

This blog series creates a small operating system in the [Rust programming language](#). Each post is a small tutorial and includes all needed code, so you can follow along if you like. The source code is also available in the corresponding [Github repository](#).

Latest post: [Integration Tests](#)

BARE BONES

A Freestanding Rust Binary

This post describes how to create a Rust executable that does not link the standard library. This makes it possible to run Rust code on the [bare metal](#) without an underlying operating system. [read more...](#)

A Minimal Rust Kernel

In this post we create a minimal 64-bit Rust kernel. We built upon the [freestanding Rust binary](#) from the previous post to create a bootable disk image, that prints something to the screen. [read more...](#)

VGA Text Mode


The [VGA text mode](#) is a simple way to print text to the screen. In this post, we create an interface that makes its usage safe and simple, by encapsulating all unsafety in a separate module and providing support for Rust's [formatting macros](#). [read more...](#)

Recent Updates

- [No more RUST_TARGET_PATH](#)
May 07
- [Rust automatically injects a dependency on `compiler_builtins` now](#) Apr 08
- [Mention RUST_TARGET_PATH in Set Up Rust post \(first edition\)](#) Mar 13
- [First bits of the second edition](#) Feb 10
- [Use proper size for heap init](#) Dec 15

OS Development in Rust

- **Writing an OS in Rust:** Tutorials for basic functionality
- **Redox OS:** Most complete Rust OS, microkernel design



OS Development in Rust

- **Writing an OS in Rust:** Tutorials for basic functionality
- **Redox OS:** Most complete Rust OS, microkernel design
- **Tock:** Operating system for embedded systems



Programmable IoT starts at the edge

An embedded operating system designed for running multiple concurrent, mutually distrustful applications on low-memory and low-power microcontrollers.



Extensible

Safely use drivers and kernel extensions from third parties



Reliable

Run processes reliably with minimal resource overhead



Low-power

Automatic low power operation

OS Development in Rust

- **Writing an OS in Rust:** Tutorials for basic functionality
- **Redox OS:** Most complete Rust OS, microkernel design
- **Tock:** Operating system for embedded systems
- **Nebulet:** Experimental WebAssembly kernel
 - WebAssembly is a binary format for executable code in web pages
 - Idea: Run wasm applications instead of native binaries
 - Wasm is sandboxed, so it can safely run in kernel address space
 - A bit slower than native code
 - But no expensive context switches or syscalls

OS Development in Rust

- **Writing an OS in Rust:** Tutorials for basic functionality
- **Redox OS:** Most complete Rust OS, microkernel design
- **Tock:** Operating system for embedded systems
- **Nebulet:** Experimental WebAssembly kernel

What does using Rust mean for OS development?



Rust means...

Memory Safety

Memory Safety

- No invalid memory accesses
 - No buffer overflows
 - No dangling pointers
 - No data races
- Guaranteed by Rust's ownership system
 - At compile time


Memory Safety

- No invalid memory accesses
 - No buffer overflows
 - No dangling pointers
 - No data races
- Guaranteed by Rust's ownership system
 - At compile time


In C:

- Array capacity is not checked on access
 - Easy to get buffer overflows
- Every `malloc` needs exactly one `free`
 - Easy to get *use-after-free* or *double-free* bugs
- Vulnerabilities caused by memory unsafety are still common

Memory Safety




Buffer Overflow Vulnerabilities in Linux




Source: https://www.cvedetails.com/product/47/Linux-Linux-Kernel.html?vendor_id=33

Memory Safety



Linux CVEs in 2018 (Jan – Apr)



Memory Safety: A Strict Compiler

- It can take some time until your program compiles
- Lifetimes can be complicated
 - “error: x does not live long enough”

However:

- “If it compiles, it usually works”
- Far less debugging
 - No data races!
- Refactoring is safe and painless



Rust means...

Encapsulating Unsafety

Encapsulating Unsafety

- Not everything can be verified at compile time
- Sometimes you need unsafe in a kernel
 - Writing to the VGA text buffer at 0xb8000
 - Modifying CPU configuration registers
 - Switching the address space (reloading CR3)

Encapsulating Unsafety

- Not everything can be verified at compile time
- Sometimes you need unsafe in a kernel
 - Writing to the VGA text buffer at 0xb8000
 - Modifying CPU configuration registers
 - Switching the address space (reloading CR3)
- Rust has unsafe blocks that allow to
 - Dereference raw pointers
 - Call unsafe functions
 - Access mutable statics
 - Implement unsafe traits
- Goal: Provide safe abstractions that encapsulate unsafety
 - Like hardware abstractions in an OS

Encapsulating Unsafety: Example

```
/// Read current page table
pub fn read_cr3() -> PhysFrame { ... }

/// Load a new page table
pub unsafe fn write_cr3(frame: PhysFrame) { ... }

/// Invalidate the TLB completely by reloading the CR3 register.
pub fn flush_tlb() { // safe interface
    let frame = read_cr3();
    unsafe { write_cr3(frame) }
}
```

- The CR3 register holds the root page table address
 - Reading is safe
 - Writing is unsafe (because it changes the address mapping)
- The `flush_tlb` function provides a safe interface
 - It can't be used in an unsafe way



Rust means...

A Powerful Type System

A Powerful Type System: Mutexes

C++

```
std::vector data = {1, 2, 3};  
// mutex is unrelated to data  
std::mutex mutex;  
  
// unsynchronized access possible  
data.push_back(4);  
  
mutex.lock();  
data.push_back(5);  
mutex.unlock();
```

Rust

```
let data = vec![1, 2, 3];  
// mutex owns data  
let mutex = Mutex::new(data);  
  
// compilation error: data was moved  
data.push(4);  
  
let mut d = mutex.lock().unwrap();  
d.push(5);  
// released at end of scope
```

⇒ Rust ensures that Mutex is locked before accessing data

A Powerful Type System: Page Table Methods

Add a page table mapping:

```
fn map_to<S: PageSize>(
    &mut PageTable,
    page: Page<S>,           // map this page
    frame: PhysFrame<S>,     // to this frame
    flags: PageTableFlags,
) {...}

impl PageSize for Size4KB {...} // standard page
impl PageSize for Size2MB {...} // “huge” 2MB page
impl PageSize for Size1GB {...} // “giant” 1GB page (only on some architectures)
```

A Powerful Type System: Page Table Methods

Add a page table mapping:

```
fn map_to<S: PageSize>(  
    &mut PageTable,  
    page: Page<S>,           // map this page  
    frame: PhysFrame<S>,     // to this frame  
    flags: PageTableFlags,  
) {...}  
  
impl PageSize for Size4KB {...} // standard page  
impl PageSize for Size2MB {...} // “huge” 2MB page  
impl PageSize for Size1GB {...} // “giant” 1GB page (only on some architectures)
```

- Generic over the page size
 - 4KB, 2MB or 1GB
- Page and frame must have the same size

A Powerful Type System

Allows to:

- Make misuse impossible
 - Impossible to access data behind a Mutex without locking
- Represent contracts in code instead of documentation
 - Page size of page and frame parameters must match in `map_to`

Everything happens at compile time \Rightarrow *No run-time cost!*



Rust means...

Easy Dependency Management

Easy Dependency Management

- Over 15000 crates on **crates.io**
- Simply specify the desired version
 - Add single line to `Cargo.toml`
- Cargo takes care of the rest
 - Downloading, building, linking



Easy Dependency Management

- Over 15000 crates on **crates.io**
 - Simply specify the desired version
 - Add single line to `Cargo.toml`
 - Cargo takes care of the rest
 - Downloading, building, linking
-
- It works the same for OS kernels
 - Crates need to be `no_std`
 - Useful crates: `bitflags`, `spin`, `arrayvec`, `x86_64`, ...



Easy Dependency Management

- **bitflags**: A macro for generating structures with single-bit flags

```
#[macro_use] extern crate bitflags;

bitflags! {
    pub struct PageTableFlags: u64 {
        const PRESENT = 1 << 0;    // bit 0
        const WRITABLE = 1 << 1;    // bit 1
        const HUGE_PAGE = 1 << 7;    // bit 7
        ...
    }
}

fn main() {
    let stack_flags = PageTableFlags::PRESENT | PageTableFlags::WRITABLE;
    assert_eq!(stack_flags.bits(), 0b11);
}
```

Easy Dependency Management

- **bitflags**: A macro for generating structures with single-bit flags
- **spin**: Spinning synchronization primitives such as spinlocks

Easy Dependency Management

- **bitflags**: A macro for generating structures with single-bit flags
- **spin**: Spinning synchronization primitives such as spinlocks
- **arrayvec**: Stack-based vectors backed by a fixed sized array

```
use arrayvec::ArrayVec;  
  
let mut vec = ArrayVec::<[i32; 16]>::new();  
vec.push(1);  
vec.push(2);  
assert_eq!(vec.len(), 2);  
assert_eq!(vec.as_slice(), &[1, 2]);
```

Easy Dependency Management

- **bitflags**: A macro for generating structures with single-bit flags
- **spin**: Spinning synchronization primitives such as spinlocks
- **arrayvec**: Stack-based vectors backed by a fixed sized array
- **x86_64**: Structures, registers, and instructions specific to x86_64
 - Control registers
 - I/O ports
 - Page Tables
 - Interrupt Descriptor Tables
 - ...

Easy Dependency Management

- **bitflags**: A macro for generating structures with single-bit flags
 - **spin**: Spinning synchronization primitives such as spinlocks
 - **arrayvec**: Stack-based vectors backed by a fixed sized array
 - **x86_64**: Structures, registers, and instructions specific to x86_64
-
- Over 350 crates in the no_std category
 - Many more can be trivially made no_std



Rust means...

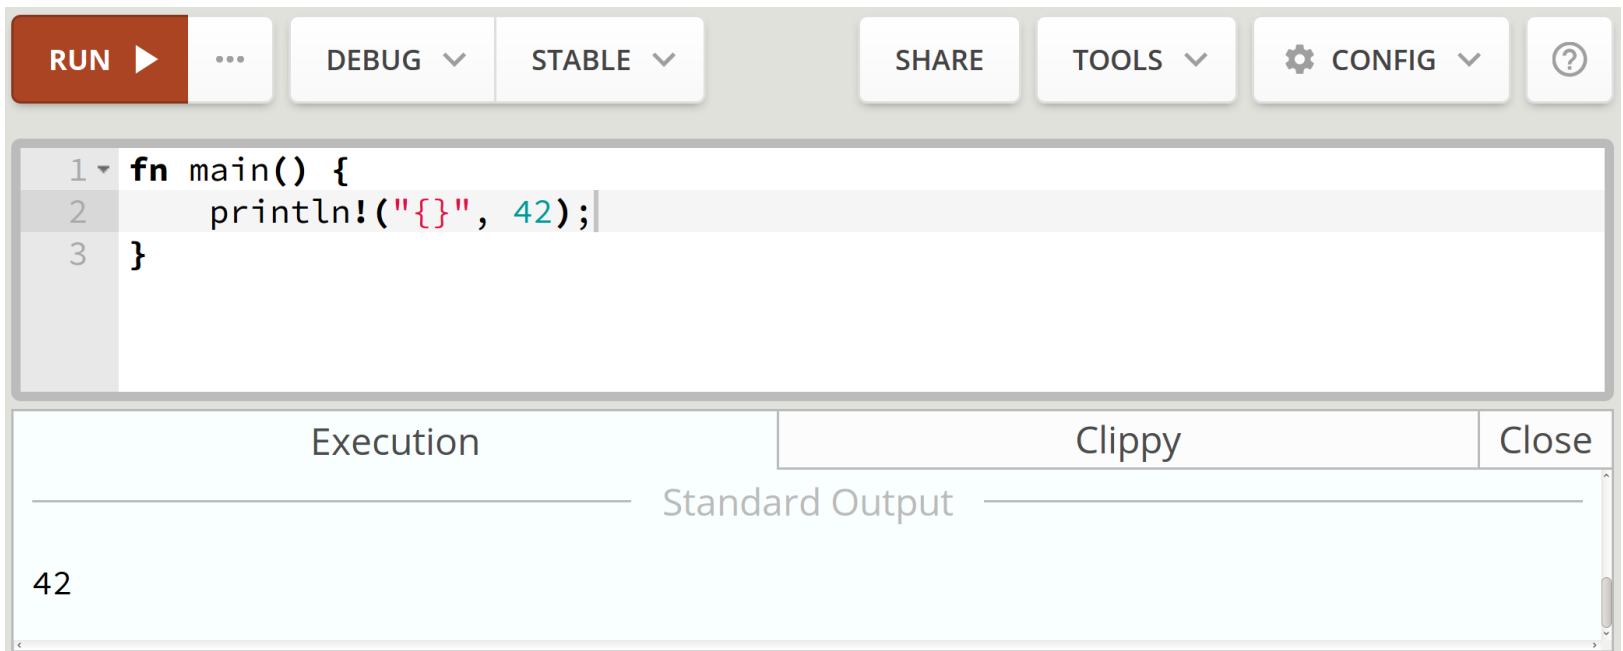
Great Tooling

Great Tooling

- **rustup**: Use multiple Rust versions for different directories
- **cargo**: Automatically download, build, and link dependencies
- **rustfmt**: Format Rust code according to style guidelines

Great Tooling

- **rustup**: Use multiple Rust versions for different directories
- **cargo**: Automatically download, build, and link dependencies
- **rustfmt**: Format Rust code according to style guidelines
- **Rust Playground**: Run and share code snippets in your browser



Great Tooling

- **clippy**: Additional warnings for dangerous or unidiomatic code

```
fn equal(x: f32, y: f32) -> bool {  
    if x == y { true } else { false }  
}
```

Great Tooling

- **clippy**: Additional warnings for dangerous or unidiomatic code

```
fn equal(x: f32, y: f32) -> bool {  
    if x == y { true } else { false }  
}
```

error: strict comparison of f32 or f64

--> src/main.rs:2:8

|

2 | if x == y { true } else { false }

| ^^^^^ help: consider comparing them within some error: (x - y).abs() < err

warning: this if-then-else expression returns a bool literal

--> src/main.rs:2:5

|

2 | if x == y { true } else { false }

| ^^^ help: you can reduce it to: x == y

Great Tooling

- **proptest**: A property testing framework

```
fn parse_date(s: &str) -> Option<(u32, u32, u32)> {  
    // [...] check if valid YYYY-MM-DD format  
    let year = &s[0..4];  
    let month = &s[6..7]; // BUG: should be 5..7  
    let day = &s[8..10];  
    convert_to_u32(year, month, date)  
}  
  
proptest! {  
    #[test]  
    fn parse_date(y in 0u32..10000, m in 1u32..13, d in 1u32..32) {  
        let date_str = format!("{:04}-{:02}-{:02}", y, m, d);  
        let (y2, m2, d2) = parse_date(&date_str).unwrap();  
        prop_assert_eq!((y, m, d), (y2, m2, d2));  
    }  
}
```


Great Tooling for OS Development

In C:

- First step is to build a **cross compiler**
 - A gcc that compiles for a bare-metal target system
 - Lots of build dependencies
- On Windows, you have to use **cygwin**
 - Required for using the GNU build tools (e.g. make)
 - The *Windows Subsystem for Linux* might also work

In Rust:

- Rust works natively on Linux, Windows, and macOS
- The Rust compiler `rustc` is already a cross-compiler
- For linking, we can use the cross-platform **lld** linker
 - By the LLVM project

Great Tooling for OS Development

bootimage: Create a bootable disk image from a Rust kernel

- Cross-platform, no C dependencies
- Automatically downloads and compiles a bootloader
- **bootloader:** A x86 bootloader written in Rust and inline assembly

Goals:

- Make building your kernel as easy as possible
- Let beginners dive immediately into OS programming
 - No hours-long toolchain setup
- Remove platform-specific differences
 - You shouldn't need Linux to do OS development

Great Tooling for OS Development

In development: **bootimage test**

- Basic integration test framework
- Runs each test executable in an isolated QEMU instance
 - Tests are completely independent
 - Results are reported through the serial port
- Allows testing in target environment

Testing on real hardware?



Rust means...

An Awesome Community

An Awesome Community

- Code of Conduct from the beginning
 - “We are committed to providing a **friendly, safe and welcoming environment** for all [...]”
 - “We will exclude you from interaction if you insult, demean or harass anyone”
 - Followed on GitHub, IRC, the Rust subreddit, etc.

An Awesome Community

- Code of Conduct from the beginning
 - “We are committed to providing a **friendly, safe and welcoming environment** for all [...]”
 - “We will exclude you from interaction if you insult, demean or harass anyone”
 - Followed on GitHub, IRC, the Rust subreddit, etc.
- It works!
 - No inappropriate comments for “Writing an OS in Rust” so far
 - Focused technical discussions

An Awesome Community

- Code of Conduct from the beginning
 - “We are committed to providing a **friendly, safe and welcoming environment** for all [...]”
 - “We will exclude you from interaction if you insult, demean or harass anyone”
 - Followed on GitHub, IRC, the Rust subreddit, etc.
- It works!
 - No inappropriate comments for “Writing an OS in Rust” so far
 - Focused technical discussions

vs:

“So this patch is utter and absolute garbage, and should be shot in the head and buried very very deep.”

[Linus Torvalds on 14 Aug 2017](#)



Rust means...

No Elitism

No Elitism

Typical elitism in OS development:

“A **decade of programming**, including a few years of low-level coding in assembly language and/or a systems language such as C, is **pretty much the minimum necessary** to even understand the topic well enough to work in it.”


From [wiki.osdev.org/Beginner Mistakes](http://wiki.osdev.org/Beginner_Mistakes)

Most Rust Projects:

- It doesn't matter where you come from
 - C, C++, Java, Python, JavaScript, ...
- It's fine to ask questions
 - People are happy to help

No Elitism

- **IntermezzOS**: “People who have not done low-level programming before are a specific target of this book”



intermezz**OS**, (a little OS)

- **Writing an OS in Rust**: Deliberately no particular target audience
 - People are able to decide themselves
 - Provide links for things not explained on the blog
 - E.g. for advanced Rust and OS concepts



Rust means...

Exciting New Features

Exciting New Features

- **Impl Trait:** Return closures from functions
- **Non-Lexical Lifetimes:** A more intelligent borrow checker
- **WebAssembly:** Run Rust in browsers

Exciting New Features

- **Impl Trait:** Return closures from functions
- **Non-Lexical Lifetimes:** A more intelligent borrow checker
- **WebAssembly:** Run Rust in browsers

In development: **Futures and async / await**

- Simple and fast asynchronous code
- How does it work?
- What does it mean for OS development?

Futures

Result of an asynchronous computation:

```
trait Future {  
    type Item;  
    type Error;  
    fn poll(&mut self, cx: &mut Context) -> Result<Async<Self::Item>,  
                                                Self::Error>;  
}  
  
enum Async<T> {  
    Ready(T),  
    Pending,  
}
```

- Instead of blocking, `Async::Pending` is returned

Futures: Implementation Details

- Futures do nothing until polled
- An Executor is used for polling multiple futures until completion
 - Like a scheduler
- If future is not ready when polled, a Waker is created
 - Notifies the Executor when the future becomes ready
 - Avoids continuous polling

Combinators

- Transform a future without polling it (similar to iterators)
- Examples
 - `future.map(|v| v + 1)`: Applies a function to the result
 - `future_a.join(future_b)`: Wait for both futures
 - `future.and_then(|v| some_future(v))`: Chain dependent futures

Async / Await

Traditional synchronous code:

```
fn get_user_from_database(user_id: u64) -> Result<User> {...}

fn handle_request(request: Request) -> Result<Response> {
    let user = get_user_from_database(request.user_id)?;
    generate_response(user)
}
```

- Thread blocked until database read finished
 - Complete thread stack unusable
- Number of threads limits number of concurrent requests

Async / Await

Asynchronous variant:

```
async fn get_user_from_database(user_id: u64) -> Result<User> {...}

async fn handle_request(request: Request) -> Result<Response> {
    let future = get_user_from_database(request.user_id);
    let user = await!(future)?;
    generate_response(user)
}
```

- Async functions return `Future<Item=T>` instead of `T`
- No blocking occurs
 - Stack can be reused for handling other requests
- Thousands of concurrent requests possible

How does `await` work?

Async / Await: Generators

- Functions that can suspend themselves via `yield`:

```
fn main() {  
    let mut generator = || {  
        println!("2");  
        yield;  
        println!("4");  
    };  
  
    println!("1");  
    unsafe { generator.resume() };  
    println!("3");  
    unsafe { generator.resume() };  
    println!("5");  
}
```

Async / Await: Generators

- Functions that can suspend themselves via `yield`:

```
fn main() {  
    let mut generator = || {  
        println!("2");  
        yield;  
        println!("4");  
    };  
  
    println!("1");  
    unsafe { generator.resume() };  
    println!("3");  
    unsafe { generator.resume() };  
    println!("5");  
}
```

- Compiled as state machines

```

let mut generator = {
    enum Generator { Start, Yield1, Done, }

    impl Generator {
        unsafe fn resume(&mut self) {
            match self {
                Generator::Start => {
                    println!("2");
                    *self = Generator::Yield1;
                }
                Generator::Yield1 => {
                    println!("4");
                    *self = Generator::Done;
                }
                Generator::Done => panic!("generator resumed after completion")
            }
        }
    }
    Generator::Start
};

```

Async / Await: Generators

- Generators can keep state:

```
fn main() {  
    let mut generator = || {  
        let number = 42;  
        let ret = "foo";  
  
        yield number; // yield can return values  
        return ret  
    };  
  
    unsafe { generator.resume() };  
    unsafe { generator.resume() };  
}
```

Where are number and ret stored between resume calls?

```

let mut generator = {
    enum Generator {
        Start(i32, &'static str),
        Yield1(&'static str),
        Done,
    }
    impl Generator {
        unsafe fn resume(&mut self) -> GeneratorState<i32, &'static str> {
            match self {
                Generator::Start(i, s) => {
                    *self = Generator::Yield1(s); GeneratorState::Yielded(i)
                }
                Generator::Yield1(s) => {
                    *self = Generator::Done; GeneratorState::Complete(s)
                }
                Generator::Done => panic!("generator resumed after completion")
            }
        }
    }
    Generator::Start(42, "foo")
};

```

Async / Await: Implementation

```
async fn handle_request(request: Request) -> Result<Response> {  
    let future = get_user_from_database(request.user_id);  
    let user = await!(future)?;  
    generate_response(user)  
}
```

Compiles roughly to:

```
async fn handle_request(request: Request) -> Result<Response> { GenFuture(|| {  
    let future = get_user_from_database(request.user_id);  
    let user = loop { match future.poll() {  
        Ok(Async::Ready(u)) => break Ok(u),  
        Ok(Async::NotReady) => yield,  
        Err(e) => break Err(e),  
    } }?;  
    generate_response(user)  
}}}
```

Async / Await: Implementation

Transform Generator into Future:

```
struct GenFuture<T>(T);

impl<T: Generator> Future for GenFuture<T> {
    fn poll(&mut self, cx: &mut Context) -> Result<Async<T::Item>, T::Error> {
        match unsafe { self.0.resume() } {
            GeneratorStatus::Complete(Ok(result)) => Ok(Async::Ready(result)),
            GeneratorStatus::Complete(Err(e)) => Err(e),
            GeneratorStatus::Yielded => Ok(Async::NotReady),
        }
    }
}
```

Async / Await: For OS Development?

- Everything happens at compile time
 - Can be used in OS kernels and on embedded devices
- Makes asynchronous code simpler

Use Case: **Cooperative multitasking**

- Yield when waiting for I/O
- Executor then polls another future
- Interrupt handler notifies Waker
- Only a single thread is needed
 - Devices with limited memory

Async / Await: An OS Without Blocking?

A blocking thread makes its whole stack unusable

- Makes threads heavy-weight
- Limits the number of threads in the system

What if blocking was not allowed?

Async / Await: An OS Without Blocking?

A blocking thread makes its whole stack unusable

- Makes threads heavy-weight
- Limits the number of threads in the system

What if blocking was not allowed?

- Threads would return futures instead of blocking
- Scheduler would schedule futures instead of threads
- Stacks could be reused for different threads
- Only a few stacks are needed for many, many futures
- Task-based instead of thread-based concurrency
 - Fine grained concurrency at the OS level

Summary

Rust means:

- **Memory Safety** no overflows, no invalid pointers, no data races
- **Encapsulating Unsafety** creating safe interfaces
- **A Powerful Type System** make misuse impossible
- **Easy Dependency Management** cargo, crates.io
- **Great Tooling** clippy, proptest, bootimage
- **An Awesome Community** code of conduct
- **No Elitism** asking questions is fine, no minimum requirements
- **Exciting New Features** futures, async / await

Slides are available at <https://os.phil-opp.com/talks>

Summary

Rust means:

- **Memory Safety** no overflows, no invalid pointers, no data races
- **Encapsulating Unsafety** creating safe interfaces
- **A Powerful Type System** make misuse impossible
- **Easy Dependency Management** cargo, crates.io
- **Great Tooling** clippy, proptest, bootimage
- **An Awesome Community** code of conduct
- **No Elitism** asking questions is fine, no minimum requirements
- **Exciting New Features** futures, async / await

Slides are available at <https://os.phil-opp.com/talks>

Extra Slides

Encapsulating Unsafety

Not possible in all cases:

```
/// Write a new root table address into the CR3 register.  
pub fn write_cr3(page_table_frame: PhysFrame, flags: Cr3Flags) {  
    let addr = page_table_frame.start_address();  
    let value = addr.as_u64() | flags.bits();  
    unsafe { asm!("mov $0, %cr3" :: "r" (value) : "memory"); }  
}
```

Encapsulating Unsafety

Not possible in all cases:

```
/// Write a new root table address into the CR3 register.  
pub fn write_cr3(page_table_frame: PhysFrame, flags: Cr3Flags) {  
    let addr = page_table_frame.start_address();  
    let value = addr.as_u64() | flags.bits();  
    unsafe { asm!("mov $0, %cr3" :: "r" (value) : "memory"); }  
}
```

Problem: Passing an invalid PhysFrame could break memory safety!

- A frame that is no page table
- A page table that maps all pages to the same frame
- A page table that maps two random pages to the same frame

Encapsulating Unsafety

Not possible in all cases:

```
/// Write a new root table address into the CR3 register.  
pub unsafe fn write_cr3(page_table_frame: PhysFrame, flags: Cr3Flags) {  
    let addr = page_table_frame.start_address();  
    let value = addr.as_u64() | flags.bits();  
    asm!("mov $0, %cr3" :: "r" (value) : "memory");  
}
```

Problem: Passing an invalid PhysFrame could break memory safety!

- A frame that is no page table
- A page table that maps all pages to the same frame
- A page table that maps two random pages to the same frame

⇒ Function needs to be **unsafe** because it depends on valid input

Encapsulating Unsafety

Edge Cases: Functions that...

- ... disable paging?

Encapsulating Unsafety

Edge Cases: Functions that...

- ... disable paging? unsafe
- ... disable CPU interrupts?

Encapsulating Unsafety

Edge Cases: Functions that...

- ... disable paging? unsafe
- ... disable CPU interrupts? safe
- ... might cause CPU exceptions?

Encapsulating Unsafety

Edge Cases: Functions that...

- ... disable paging? unsafe
- ... disable CPU interrupts? safe
- ... might cause CPU exceptions? safe
- ... can be only called from privileged mode?

Encapsulating Unsafety

Edge Cases: Functions that...

- ... disable paging? unsafe
- ... disable CPU interrupts? safe
- ... might cause CPU exceptions? safe
- ... can be only called from privileged mode? safe
- ... assume certain things about the hardware?
 - E.g. there is a VGA text buffer at 0xb8000

Encapsulating Unsafety

Edge Cases: Functions that...

- ... disable paging? unsafe
- ... disable CPU interrupts? safe
- ... might cause CPU exceptions? safe
- ... can be only called from privileged mode? safe
- ... assume certain things about the hardware? depends
 - E.g. there is a VGA text buffer at 0xb8000

Async / Await: Generators

- Why is resume unsafe?

```
fn main() {  
    let mut generator = move || {  
        let foo = 42;  
        let bar = &foo;  
        yield;  
        return bar  
    };  
    unsafe { generator.resume() };  
    let heap_generator = Box::new(generator);  
    unsafe { heap_generator.resume() };  
}
```

```
enum Generator {  
    Start,  
    Yield1(i32, &i32),  
    Done,  
}
```

Async / Await: Generators

- Why is resume unsafe?

```
fn main() {  
    let mut generator = move || {  
        let foo = 42;  
        let bar = &foo;  
        yield;  
        return bar  
    };  
    unsafe { generator.resume() };  
    let heap_generator = Box::new(generator);  
    unsafe { heap_generator.resume() };  
}
```

```
enum Generator {  
    Start,  
    Yield1(i32, &i32),  
    Done,  
}
```

- Generator contains reference to itself
 - No longer valid when moved to the heap \Rightarrow undefined behavior
 - Must not be moved after first resume

Await: Just Syntactic Sugar?

Is `await` just syntactic sugar for the `and_then` combinator?

```
async fn handle_request(request: Request) -> Result<Response> {  
    let user = await!(get_user_from_database(request.user_id))?;  
    generate_response(user)  
}
```

```
async fn handle_request(request: Request) -> Result<Response> {  
    get_user_from_database(request.user_id).and_then(|user| {  
        generate_response(user)  
    })  
}
```

In this case, both variants work.

Await: Not Just Syntactic Sugar!

```
fn read_info_buf(socket: &mut Socket) -> [u8; 1024]
    -> impl Future<Item = [0; 1024], Error = io::Error> + 'static
{
    let mut buf = [0; 1024];
    let mut cursor = 0;
    while cursor < 1024 {
        cursor += await!(socket.read(&mut buf[cursor..]))?;
    };
    buf
}
```

- We don't know how many `and_then` we need
 - But each one is their own type -> boxed trait objects required
- `buf` is a local stack variable, but the returned future is `'static`
 - Not possible with `and_then`
 - *Pinned types* allow it for `await`